

# Experiment Design, Group 04 – Reproducibility - Paper 2

Helmuth Breitenfellner  
e8725866@student.tuwien.ac.at

László Király  
e9227679@student.tuwien.ac.at

Gerald Weber  
e0125536@student.tuwien.ac.at

## Abstract

This documents the activities and results performed by Group 04 during *Exercise 2: Reproducibility* of the lecture *188.992 Experiment Design for Data Science*.

## A) Paper selection

We selected Option 2 - Prediction of music genre across different taxonomies which is based on the paper *MediaEval 2018 AcousticBrainz Genre Task: A Baseline Combining Deep Feature Embeddings Across Datasets* written by Sergio Oramas, Dmitry Bogdanov and Alastair Porter for the MediaEval conference 2018 in France. We have chosen the paper because we were curious to reproduce the output of the Neural Network and see what are the problems here in difference to regular Machine Learning papers.

### (1) Paper Introduction

The paper is a baseline approach for the MediaEval 2018 AcousticBrainz Genre Task based on a deep neural network. *The task is focused on content-based musicgenre recognition using genre annotations from multiple sources and large-scale music features data available in the AcousticBrainz database* [BPUS17].

### (2) Datasets

The conference provides a website for the classification task<sup>1</sup> which describes the task, the schedule and the dataset on a subpage<sup>2</sup> in detail. This subpage contains multiple links:

- Zenodo<sup>3</sup> uploaded by Dmitry Bogdanov, Alastair Porter et al.
- Google Drive<sup>4</sup> containing train/test/validation folders provided by Dmitry Bogdanov and Alastair Porter

The first Zenodo page contains all the required data (lastfm, tag) for the classification task. The second Zenodo page contains restricted access to the AllMusic ground truth, which has been granted after a simple request on the Zenodo page. The compressed file size is about 47GB.

After decompressing the dataset is structured as followed:

z-mediaeval-train 83GB, 1.458.447 files

diaeval-validation 18GB, 313.860 files

- acousticbrainz-mediaeval-allmusic-train.tsv 260MB (restricted dataset)
- acousticbrainz-mediaeval-allmusic-validation.tsv 55MB (restricted dataset)
- acousticbrainz-mediaeval-discogs-train.tsv 127MB
- acousticbrainz-mediaeval-discogs-validation.tsv 26MB
- acousticbrainz-mediaeval-lastfm-train.tsv 62MB
- acousticbrainz-mediaeval-lastfm-validation.tsv 14MB
- acousticbrainz-mediaeval-tagtraum-train.tsv 59MB

- acousticbrainz-mediaeval-tagtraum-validation.tsv 13MB

Total size of the extracted files: 102GB

The folders contains subfolders ranging from 00-ff containing JSON files. These files contains precomputed Audio features extracted with Essentia<sup>5</sup> from the community-built AcousticBrainz database. The .tsv files contains features in the format: recordingmbid releasegroupmbid genre1 genre2 genre3 genre4 genre5 genre6 genre7 genre8 genre9 genre10 genre11 genre12 genre13 genre14 genre15 genre16 genre17 genre18 genre19 genre20 genre21 genre22 genre23 genre24 genre25 genre26 genre27 genre28 genre29 genre30 genre31 genre32 genre33 genre34 genre35 genre36 genre37 genre38 genre39 genre40

The field *recordingmbid* is the MusicBrainz identifier of the particular recording, whereas *releasegroupmbid* is a MusicBrainz identifier of a release group (an album, single, or compilation) that it belongs to. The related genres and subgenres are encoded into the columns *genre1-genre40*. For each recordingmbid there exists a file the *acousticbrainz-mediaeval-train* folder.

## B) Paper - Reproduction

The authors of the paper provided a link to a Github repository<sup>6</sup> which contains a data-preparation folder and a link to another repository<sup>7</sup>. This linked repository contains three experiments written by Oramas et al. . Based on the description of the experiments and on the number of features described in the paper (2669) we could identify experiment 3: *TISMIR Experiments (Singlelabel Classification)*<sup>8</sup> as the baseline approach.

With this information we inspected the *datapreparation/create\_h5.py* file to find out how to prepare the data. The existing file *datapreparation/create\_h5.py* requires some *features.clean.std.csv* and *.clean.std.csv* as well as *genres.csv* as input. None of them is available in the dataset.

Tries to contact the main author via two mail addresses (soram@pandora.com, sergio.oramas@upf.edu) were not successful.

## C) Reproducing #2 - A new hope

Some online research led us to a recently created Github repository<sup>9</sup> which deals with the same paper. Additionally it has a preprocessing step included. Naturally, we cloned the repository and started

<sup>1</sup>, seen on 2020-01-30

<sup>2</sup>, seen on 2020-01-30

<sup>3</sup>, seen on 2020-01-30

<sup>4</sup>

<sup>5</sup>

<sup>6</sup>, seen on 2020-01-30

<sup>7</sup>, seen on 2020-01-30

<sup>8</sup>denoted in the *run\_experiments.py* as *ISMIR 2019 Experiments (Baseline for AcousticBrainz Genre Dataset Classification)*

<sup>9</sup>, seen on 2020-01-30

**(1) Datasets**

**(2) Sources**

**(3) Doing**

**References**

[BPUS17] Dmitry Bogdanov, Alastair Porter, Julián Urbano, and Hendrik Schreiber.  
The mediaeval 2017 acousticbrainz genre task: content-based music genre

recognition from multiple sources. 01 2017.